

Prediction ratio vs. measured latency, all kernels
(dashed line = perfect prediction)

